

Note - eq 3.47 One line isn't so clear, but

$$\begin{aligned}
& (VD^2V^T + \lambda I)(VD^2V^T + \lambda I)^{-1} = I \\
\Rightarrow & (VD^2V^T + \lambda I)(VD^2V^T + \lambda I)^{-1}V = V \\
\Rightarrow & (D^2V^T + \lambda V^T)(VD^2V^T + \lambda I)^{-1}V = I \\
\Rightarrow & (D^2 + \lambda I)V^T(VD^2V^T + \lambda I)^{-1}V = I
\end{aligned}$$

and so

$$V^T(VD^2V^T + \lambda I)^{-1}V = (D^2 + \lambda I)^{-1}.$$

3.1 Recall the method of using repeated simple linear regression to do multiple linear regression: We form an orthogonal spanning set z_1, \dots, z_p of the column space in such a way that z_1, \dots, z_k spans the span of the first k columns of X . We finally regress y onto z_p to obtain $\hat{\beta}_p$. But we can do more: If $\gamma_j = \langle y, z_j \rangle$, then $\hat{y} = \sum_j \frac{\gamma_j}{\|z_j\|^2} z_j$ is the orthogonal projection onto the column space of X in the full rank case. What's more, $\sum_{j < p} \frac{\gamma_j}{\|z_j\|^2} z_j$ is the result of using only the first $p - 1$ features. Let $r_p = y - \hat{y}$. Recall that $\hat{\beta}_p = \frac{\gamma_p}{\|z_p\|^2}$, and so

$$\begin{aligned}
\text{RSS}_0 &= \left\| y - \sum_{j < p} \frac{\gamma_j}{\|z_j\|^2} z_j \right\|^2 \\
&= \left\| r_p + \hat{\beta}_p z_p \right\|^2 \\
&= \|r_p\|^2 + \hat{\beta}_p^2 \|z_p\|^2 \\
\text{RSS}_1 &= \|y - \hat{y}\|^2 \\
&= \|r_p\|^2
\end{aligned}$$

So we can express the formula for the F-score

$$\begin{aligned}
F &= \frac{\text{RSS}_0 - \text{RSS}_1}{\text{RSS}_1 / (N - p - 1)} \\
&= \frac{\hat{\beta}_p^2 \|z_p\|^2}{\hat{\sigma}^2}
\end{aligned}$$

Now, write $X = Z\Gamma$, where Z 's columns are the z_i 's and Γ is upper triangular with diagonal entries 1. Then $(X^T X)^{-1} = (\Gamma^T Z^T Z \Gamma)^{-1} = (\Gamma^T D \Gamma)^{-1}$. Here, D is diagonal on $(\|z_1\|^2, \dots, \|z_p\|^2)$. Since we assume that $X^T X$ is invertible, both Γ and D are, so $(\Gamma^T D \Gamma)^{-1} = \Gamma^{-1} D^{-1} \Gamma^{-T}$. Now, we can compute the

lower-right entry

$$\begin{aligned}
(\Gamma^{-1} D^{-1} \Gamma^{-T})_{pp} &= \sum_{ij} (\Gamma^{-1})_{ip} D_{ij}^{-1} (\Gamma^{-T})_{jp} \\
&= \sum_i \frac{1}{\|z_i\|^2} (\Gamma^{-1})_{ip} (\Gamma^{-T})_{ip} \\
&= \frac{1}{\|z_p\|^2} (\Gamma_{pp}^{-1})^2 \\
&= \frac{1}{\|z_p\|^2}
\end{aligned}$$

and so we finally find that the F -score can be written as

$$\frac{\hat{\beta}_p^2}{\hat{\sigma}^2 d_p}$$

where d_p is the p th diagonal entry of $(X^T X)^{-1}$. By relabelling the features, we find that this also works for any of the features, and so the F -score is the square of the Z -score.

2 I would expect that the pointwise confidence intervals would be narrower. Here is a heuristic argument: Suppose that β_1, β_2 are normally distributed with mean 0 and covariance I . Then one choice of a 95% confidence set is a circle with radius $R_{95} = \log(20) \approx 3$ about the origin. On the other hand, a 95% confidence set for just β_0 is an interval of radius about 1.64, so the 2D set will have much more extreme values of β_0 . In particular, the band for $x_0 = 0$ will be much wider for the 2D confidence region. Another way to see it: If I generate samples for each point individually, I will implicitly be drawing samples from many, many more functions when I make the final plot, and so will have a narrower gap.

Simulations in the associated notebook.

3.3.a The variance of $c^T y$ is

$$\begin{aligned}
&\text{Var}\left(\sum_i c_i y_i\right) \\
&= \sum_{ij} c_i c_j \text{Cov}(y_i, y_j) \\
&= \sum_i c_i^2 \sigma^2 \\
&= \sigma^2 \|c\|^2
\end{aligned}$$

while the fact that the estimator is unbiased is expressed exactly by the equation $(X^T c)^T \beta = \alpha^T \beta$. Now, I'm going to make a small assumption. We know that the estimator is not supposed to depend on β , which is unobservable, and I'm

going to interpret this mathematically by taking $(X^T c)^T \beta = \alpha^T \beta$ to hold for *all* β . (This is certainly true of the OLS estimator, and it seems that an estimator that did not satisfy this property would be rather useless.)

Write $c = kX(X^T X)^{-1}\alpha + v$ for some constant k and some v orthogonal to $X(X^T X)^{-1}\alpha$, we find $X^T c = k\alpha + X^T v$ and taking $\beta = (X^T X)^{-1}\alpha$ gives

$$\begin{aligned}\beta^T(X^T c) &= ((X^T X)^{-1}\alpha)^T(k\alpha + X^T v) \\ &= k\alpha^T(X^T X)^{-1}\alpha + \alpha^T(X^T X)^{-1}X^T v \\ &= k\alpha^T(X^T X)^{-1}\alpha\end{aligned}$$

on one hand, and on the other we know that

$$\begin{aligned}\beta^T(X^T c) &= \beta^T \alpha \\ &= \alpha^T(X^T X)^{-1}\alpha\end{aligned}$$

and so $k = 1$. This means that

$$\sigma^2\|c\|^2 = \sigma^2\|X(X^T X)^{-1}\alpha\|^2 + \sigma^2\|v\|^2 \geq \sigma^2\alpha^T(X^T X)^{-1}\alpha,$$

which is the variance of the OLS estimator for $\alpha^T \beta$.

3.3.b A matrix A is positive-semidefinite iff $x^T A x \geq 0$ for all $x \in \mathbb{R}^n$. If $a \in \mathbb{R}^n$, then $a^T(\tilde{V} - \hat{V})a = a^T \tilde{V} a - a^T \hat{V} a$. Since \tilde{V}, \hat{V} are the covariance matrices of $\tilde{\beta}$ and $\hat{\beta}$, respectively, we find that

$$\begin{aligned}a^T \tilde{V} a &= \text{Var}(a^T \tilde{\beta}) \\ a^T \hat{V} a &= \text{Var}(a^T \hat{\beta}).\end{aligned}$$

So by the above result, $\tilde{V} - \hat{V}$ is positive semidefinite.

3.4 In the full-rank case, a single pass of the Gram-Schmidt procedure expresses $X = QR$ where Q is orthogonal and R is square and upper triangular with 1's along the diagonal. The formula

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

becomes

$$\begin{aligned}(R^T Q^T Q R)^{-1} R^T Q^T y & \\ &= (R^T R)^{-1} R^T Q^T y \\ &= R^{-1} R^{-T} R^T Q^T y \\ &= R^{-1} Q^T y\end{aligned}$$

We can compute R^{-1} during the Gram-Schmidt process as follows: Let Q_k, R_k denote the matrices formed by taking the first k columns of Q, R . We will also store the matrix R_k^{-1} . Given this data, the next step in the Gram-Schmidt process yields Q_{k+1}, R_{k+1} . We update R_{k+1}^{-1} as follows: If

$$R_{k+1} = \left(\begin{array}{c|c} R_k & Z \\ \hline 0 & a_k \end{array} \right),$$

then

$$R_{k+1}^{-1} = \left(\begin{array}{c|c} R_k^{-1} & -a_k R_k^{-1} Z \\ \hline 0 & a_k^{-1} \end{array} \right)$$

So we can store the coefficients of $\hat{\beta}$ in a list that is updated with each newly discovered column of R^{-1} and row of Q^T .

3.5 The original ridge objective is

$$\sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^p x_{ij} \beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2.$$

Rearranging it, we get

$$\sum_{i=1}^N \left(y_i - [\beta_0 + \sum_{j=1}^p \bar{x}_j \beta_j] + \sum_{j=1}^p (x_{ij} - \bar{x}_j) \beta_j \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$

So the minimum $\hat{\beta}_\lambda^{ridge}$ to this will give a minimum to

$$\sum_{i=1}^N \left(y_i - \beta_0^c + \sum_{j=1}^p (x_{ij} - \bar{x}_j) \beta_j^c \right)^2 + \lambda \sum_{j=1}^p (\beta_j^c)^2,$$

where we simply set $\beta_0^c := \beta_0 + \sum_{j=1}^p \bar{x}_j \beta_j$ and $\beta_j^c := \hat{\beta}_j$ if $j \geq 1$. The lasso condition is exactly the same. In both cases, the objective function is differentiable with respect to β_0 , and the derivative in the lasso and ridge case is

$$\begin{aligned} & \sum_{i=1}^n 2(y_i - \beta_0^c + \sum_{j=1}^p (x_{ij} - \bar{x}_j) \beta_j^c) \\ &= N\bar{y} - N\beta_0^c, \end{aligned}$$

and so $\beta_0^c = \bar{y}$ at the minimum in each case.

3.6 Let f_B be $N(0, \tau I)$. Then

$$f_{B|Y}(\beta|y) = \frac{f_{BY}(\beta, y)}{f_Y(y)} = \frac{f_{Y|B}(y|\beta) f_B(\beta)}{f_Y(y)}.$$

We are working under the assumption that y is Gaussian noise added onto $X\beta$, so

$$f_{Y|B}(y|\beta) = \frac{1}{(2\pi\sigma)^{\frac{p}{2}}} e^{-\frac{1}{2\sigma^2} \|y - X\beta\|^2},$$

while

$$f_B(\beta) = \frac{1}{(2\pi\tau)^{\frac{p}{2}}} e^{-\frac{1}{2\tau^2} \sum_{j=1}^p \beta_j^2}.$$

and so

$$f_{Y|B}(y|\beta)f_B(\beta) = \frac{1}{2\pi(\tau\sigma)^{\frac{p}{2}}} \exp \left\{ \frac{-1}{2} \left(\frac{\|y - X\beta\|^2}{\sigma^2} + \frac{\|\beta\|^2}{\tau^2} \right) \right\}.$$

Since this is quadratic in β , we find that this posterior distribution is also Gaussian, so the median and mode match. We'll find the value by taking differentiating with respect to β_K :

$$f_{Y|B}(y|\beta)f_B(\beta) \left[\frac{1}{\sigma^2} \sum_i (y_i - \sum_j x_{ij}\beta_j)(x_{iK}) - \frac{1}{\tau^2} \beta_K \right]$$

Setting this equal to zero for each component gives the vector equation

$$\frac{1}{\sigma^2} X^T (y - X\beta) = \frac{1}{\tau^2} \beta,$$

and solving that yields

$$\beta = (X^T X + \frac{\sigma^2}{\tau^2} I)^{-1} X^T y,$$

which is the ridge regression solution for which $\lambda = \frac{\sigma^2}{\tau^2}$.

3.7*** I'm not sure what this problem is asking about. If it's the same scenario as the previous problem, we reuse that work to find that the posterior density is

$$\frac{C}{2\pi(\tau\sigma)^{\frac{p}{2}}} \exp \left\{ \frac{-1}{2} \left(\frac{\|y - X\beta\|^2}{\sigma^2} + \frac{\|\beta\|^2}{\tau^2} \right) \right\},$$

with C chosen to make the expression integrate to 1 as a function of y . We take the log and get something that's proportional to the desired quantity, but with an extra constant added.

3.8 Let $\mathbf{1}, X_1, \dots, X_p$ denote the columns of the unnormalized matrix X . The matrix Q in the QR-decomposition can be found via the Gram-Schmidt process. Since the first column of X is all 1's, the GS process starts by making the first column of Q $\frac{1}{\sqrt{N}}\mathbf{1}$. The other columns are orthogonalized with respect to this vector:

$$X_i \mapsto \tilde{X}_i = X_i - \langle X_i, \frac{1}{\sqrt{N}}\mathbf{1} \rangle \frac{1}{\sqrt{N}}\mathbf{1} = X_i - \overline{X_i}\mathbf{1}.$$

Since the remaining columns of Q are formed by finding an orthonormal basis of the span of these vectors, which are none other than \tilde{X} 's columns, and the columns of the matrix U also form an orthogonal basis of that span, we know that the span of Q_2 is the same as the span of U .

When will the two match? If they match, we can express $\tilde{X} = Q_2 R_2$ and also $\tilde{X} = U D V^T$. Assuming $Q_2 = U$, multiplying both of these expressions by U^T will yield $R_2 = D V^T$, which means that $V = R_2^T D^{-1}$ is a lower triangular orthonormal matrix, which means it's diagonal with entries ± 1 . This, in turn, means that $R_2 = V D$ is diagonal, hence the columns of \tilde{X} are already orthogonal and simply need to be rescaled to become orthonormal.

Conversely, if the columns of \tilde{X} are orthogonal, then scaling each of them to have unit norm gives an orthonormal matrix U . Letting D be the diagonal matrix whose entries are the lengths of the columns of \tilde{X} , we have $\tilde{X} = U D$ is both the QR and SVD decomposition of \tilde{X} , so the first matrix in each decomposition is the same.